

churn_bank

April 19, 2023

Contexto:

Somos un banco que dispone de una base de datos con una gran cantidad de información sobre nuestros clientes. Nuestro objetivo es ayudar a los analistas a predecir la tasa de abandono de estos clientes para así poder reducirla. La base de datos incluye información demográfica como la edad, el sexo, el estado civil y la categoría de ingresos. También contiene información sobre el tipo de tarjeta, el número de meses en cartera y los periodos inactivos. Además, dispone de datos clave sobre el comportamiento de gasto de los clientes que se acercan a su decisión de cancelación. Entre esta última información hay el saldo total renovable, el límite de crédito, la tasa media de apertura a la compra y métricas analizables como el importe total del cambio del cuarto trimestre al primero o el índice medio de utilización.

Frente a este conjunto de datos podemos capturar información actualizada que puede determinar la estabilidad de la cuenta a largo plazo o su salida inminente.

Dataset:

CLIENTNUM: Identificador único para cada cliente. (Integer) Attrition_Flag: Indicador de si el cliente ha abandonado el banco o se queda (Boolean) Attrited Customer -> 0 Existing Customer -> 1 Customer_Age: Edad del cliente. (Integer) Gender: Sexo del cliente. (String) Dependent_count: Número de personas a cargo que tiene el cliente. (Integer) Education_Level: Nivel educativo del cliente. (String) Marital_Status: Marital status of customer. (String) Income_Category: Categoría de ingresos del cliente. (String) Card_Category: Tipo de tarjeta del cliente. (String) Months_on_book: El tiempo que el cliente ha estado en los libros. (Integer) Total_Relationship_Count: Número total de relaciones que tiene el cliente con el proveedor de la tarjeta de crédito. (Integer) Months_Inactive_12_mon: Número de meses que el cliente ha estado inactivo en los últimos doce meses. (Integer) Contacts_Count_12_mon: Número de contactos que ha tenido el cliente en los últimos doce meses. (Integer) Credit_Limit: Límite de crédito del cliente. (Integer) Total_Revolving_Bal: Saldo renovable total del cliente. (Integer) Avg_Open_To_Buy: Ratio medio de apertura a la compra del cliente. (Integer) Total_Amt_Chng_Q4_Q1: Importe total cambiado del trimestre 4 al trimestre 1. (Integer) Total_Trans_Amt: Importe total de la transacción. (Integer) Total_Trans_Ct: Recuento total de transacciones. (Integer) Total_Ct_Chng_Q4_Q1: Recuento total cambiado del trimestre 4 al trimestre 1. (Integer) Avg_Utilization_Ratio: Ratio de utilización media del cliente. (Integer) Months_Inactive_12_mon: Número de meses que el cliente ha estado inactivo en los últimos doce meses. (Integer) Contacts_Count_12_mon: Número de contactos que ha tenido el cliente en los últimos doce meses. (Integer) Credit_Limit: Límite de crédito del cliente. (Integer) Total_Revolving_Bal: Saldo rotativo total del cliente. (Integer) Avg_Open_To_Buy: Ratio medio de apertura a compra del cliente. (Integer) Total_Amt_Chng_Q4_Q1: Importe total cambiado del trimestre 4 al trimestre 1. (Integer) Total_Trans_Amt: Importe total de la transacción. (In-

teger) Total_Trans_Ct: Recuento total de transacciones.. (Integer) Total_Ct_Chng_Q4_Q1: Recuento total cambiado del trimestre 4 al trimestre 1. (Integer) Avg_Utilization_Ratio: Ratio de utilización media del cliente. (Integer)

1 Libraries

```
[1]: #uncomment to install libraries in google colab  
#!pip install category_encoders  
#!pip install -U imbalanced-learn
```

```
[2]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
plt.style.use('ggplot')  
%matplotlib inline  
import math  
import pickle  
from xgboost import XGBClassifier  
from xgboost import plot_importance  
import category_encoders as ce  
from sklearn.preprocessing import StandardScaler,RobustScaler  
from imblearn.over_sampling import SMOTE  
from imblearn.under_sampling import RandomUnderSampler  
from imblearn.pipeline import Pipeline  
from sklearn.model_selection import (train_test_split,  
                                     StratifiedKFold,  
                                     GridSearchCV)  
  
from sklearn.linear_model import LogisticRegression,SGDClassifier  
from sklearn.metrics import (mean_squared_error,  
                             confusion_matrix,roc_curve,classification_report,  
                             roc_auc_score,  
                             accuracy_score,f1_score)  
  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.svm import SVC  
import json
```

2 Functions

```
[3]: #1 get the number of rows for the plot  
def plot_template(n_cols,features):  
    n_rows = math.ceil(len(features)/n_cols)  
    return n_rows
```

```
[4]: #2 get the relation between two features
def attrited_feature(df,target,feature):
    attr = df[[target,feature]].groupby([feature]).mean().sort_values(by=target)
    return attr
```

```
[5]: #3 plot the attrition value in relation with a feature
def plot_target_attrition(df,col):
    plt.figure(figsize=(8,5))
    sns.barplot(data=df,x=df.index,y=col)
    plt.tight_layout()
    plt.xticks(rotation=90)
    plt.title("Attrition by {}".format(df.index.name))
    plt.show()
```

```
[6]: #4 plot confusion matrix and roc

def
↳subplots_ROC_CM(main_title,rtitle,mtitle,mod,accuracy,auc,matrix,model,Xtest,ytest,label):
↳
    fig , axes = plt.subplots(1,2,figsize=(10,8))

    y_pred_prob1_rs = mod.predict_proba(Xtest)[:,-1]
    fpr1_rs, tpr1_rs, thresholds1_rs = roc_curve(ytest, y_pred_prob1_rs)
    axes[0].plot([0, 1], [0, 1], 'k--')
    axes[0].plot(fpr1_rs, tpr1_rs, label=label)
    axes[0].legend(loc="best")
    axes[0].set_xlabel('False Positive Rate')
    axes[0].set_ylabel('True Positive Rate')
    axes[0].set_title(rtitle + "\n" +
        "AUC : {0:.2%}".format(accuracy.get(model)))

    group_counts = ["{0:0.0f}".format(v) for v in matrix.flatten()]
    group_percentages = ["{0:.2%}".format(value)
        for value in matrix.flatten()/np.sum(matrix)]
    labels = [f"{ant1}\n{ant2}" for ant1, ant2 in
↳zip(group_counts,group_percentages)]
    labels = np.asarray(labels).reshape(2,2)
    ax = sns.heatmap(matrix, annot=labels, fmt='',
        cmap='Blues', ax=axes[1])
    axes[1].set_title(mtitle + "\n" + "Accuracy Score: {0:.2%}".format(auc.
↳get(model)))
    axes[1].set_xlabel('\nPredicted Values')
    axes[1].set_ylabel('Actual Values ')
    axes[1].xaxis.set_ticklabels(['False', 'True'])
```

```
axes[1].yaxis.set_ticklabels(['False', 'True'])

plt.suptitle(main_title, size=17, weight='bold')
plt.show()
```

[7]: #5 get the item from the dict

```
def get_value(item):
    return item[1]
```

3 Load Data

[8]: `data_train = pd.read_csv("https://storage.googleapis.com/challenges_events/\n03_2023/Pre-Selection%20J0Barcelona/Data/supply_chain_train.\n↵csv", index_col="train_idx")`

`data_train.head(10).T`

```
[8]: train_idx      0      1      2  \
CLIENTNUM      713071383  714246333  718206783
Customer_Age      54      58      45
Gender            F      F      F
Dependent_count    1      4      4
Education_Level   Unknown High School Unknown
Marital_Status     Single    Married    Single
Income_Category   Unknown    Unknown Less than $40K
Card_Category     Blue      Blue      Gold
Months_on_book    36      48      36
Total_Relationship_Count    1      1      6
Months_Inactive_12_mon    3      4      1
Contacts_Count_12_mon    3      3      3
Credit_Limit    3723.0    5396.0    15987.0
Total_Revolving_Bal    1728    1803    1648
Avg_Open_To_Buy    1995.0    3593.0    14339.0
Total_Amt_Chng_Q4_Q1    0.595    0.493    0.732
Total_Trans_Amt    8554    2107    1436
Total_Trans_Ct      99      39      36
Total_Ct_Chng_Q4_Q1    0.678    0.393    1.25
Avg_Utilization_Ratio    0.464    0.334    0.103
Attrition_Flag      1      0      1

train_idx      3      4      5  \
CLIENTNUM      721096983  720028683  778942233
Customer_Age      34      49      60
Gender            F      F      F
Dependent_count    2      2      0
```

Education_Level	Graduate	High School	Doctorate
Marital_Status	Single	Married	Married
Income_Category	Less than \$40K	\$40K - \$60K	Less than \$40K
Card_Category	Blue	Blue	Blue
Months_on_book	36	39	45
Total_Relationship_Count	4	5	5
Months_Inactive_12_mon	3	3	2
Contacts_Count_12_mon	4	4	4
Credit_Limit	3625.0	2720.0	1438.3
Total_Revolving_Bal	2517	1926	648
Avg_Open_To_Buy	1108.0	794.0	790.3
Total_Amt_Chng_Q4_Q1	1.158	0.602	0.477
Total_Trans_Amt	2616	3806	1267
Total_Trans_Ct	46	61	27
Total_Ct_Chng_Q4_Q1	1.3	0.794	1.077
Avg_Utilization_Ratio	0.694	0.708	0.451
Attrition_Flag	1	1	1

train_idx	6	7	8 \
CLIENTNUM	708682908	720670458	719952408
Customer_Age	43	52	30
Gender	F	F	M
Dependent_count	4	2	0
Education_Level	Unknown	Unknown	Graduate
Marital_Status	Single	Single	Married
Income_Category	Unknown	\$40K - \$60K	Less than \$40K
Card_Category	Blue	Blue	Blue
Months_on_book	28	45	36
Total_Relationship_Count	2	3	3
Months_Inactive_12_mon	2	1	3
Contacts_Count_12_mon	1	3	2
Credit_Limit	2838.0	3476.0	2550.0
Total_Revolving_Bal	1934	1560	1623
Avg_Open_To_Buy	904.0	1916.0	927.0
Total_Amt_Chng_Q4_Q1	0.873	0.894	0.65
Total_Trans_Amt	8644	3496	1870
Total_Trans_Ct	87	58	51
Total_Ct_Chng_Q4_Q1	0.554	0.871	0.275
Avg_Utilization_Ratio	0.681	0.449	0.636
Attrition_Flag	1	1	1

train_idx	9
CLIENTNUM	708412758
Customer_Age	33
Gender	F
Dependent_count	3
Education_Level	Graduate

Marital_Status	Single
Income_Category	Less than \$40K
Card_Category	Blue
Months_on_book	36
Total_Relationship_Count	5
Months_Inactive_12_mon	2
Contacts_Count_12_mon	3
Credit_Limit	1457.0
Total_Revolving_Bal	0
Avg_Open_To_Buy	1457.0
Total_Amt_Chng_Q4_Q1	0.677
Total_Trans_Amt	2200
Total_Trans_Ct	45
Total_Ct_Chng_Q4_Q1	0.364
Avg_Utilization_Ratio	0.0
Attrition_Flag	0

```
[9]: data = data_train.copy()
```

```
[10]: test = pd.read_csv("https://storage.googleapis.com/challenges_events/\
03_2023/Pre-Selection%20JOBBarcelona/Data/supply_chain_test.\
↪csv",index_col="test_idx")
```

```
test
```

```
[10]: CLIENTNUM Customer_Age Gender Dependent_count Education_Level \
```

test_idx	CLIENTNUM	Customer_Age	Gender	Dependent_count	Education_Level
0	719455083	48	F	3	Uneducated
1	773503308	59	M	1	Uneducated
2	715452408	37	F	2	Graduate
3	711264033	47	M	3	Doctorate
4	718943508	42	M	3	Unknown
...
2021	814776033	34	M	2	Graduate
2022	720444408	35	F	1	College
2023	720503508	44	F	1	Uneducated
2024	721217283	27	M	0	Graduate
2025	770920908	39	F	0	Unknown

test_idx	Marital_Status	Income_Category	Card_Category	Months_on_book
0	Single	Less than \$40K	Blue	39
1	Single	Less than \$40K	Blue	53
2	Divorced	Less than \$40K	Blue	36
3	Divorced	\$40K - \$60K	Blue	36
4	Single	\$80K - \$120K	Blue	33
...

2021	Single	\$80K - \$120K	Blue	29
2022	Single	Less than \$40K	Blue	25
2023	Divorced	Less than \$40K	Blue	37
2024	Single	\$120K +	Blue	17
2025	Single	\$40K - \$60K	Silver	26

	Total_Relationship_Count	Months_Inactive_12_mon	\
test_idx			
0	4	3	
1	5	5	
2	4	3	
3	4	2	
4	3	3	
...	
2021	3	1	
2022	2	2	
2023	1	2	
2024	6	2	
2025	1	1	

	Contacts_Count_12_mon	Credit_Limit	Total_Revolving_Bal	\
test_idx				
0	4	2991.0	1508	
1	4	2192.0	1569	
2	3	1734.0	987	
3	3	4786.0	1516	
4	2	3714.0	2170	
...	
2021	3	13395.0	1678	
2022	4	2231.0	1791	
2023	3	5594.0	1235	
2024	2	8713.0	1354	
2025	1	22054.0	1146	

	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	\
test_idx				
0	1483.0	0.703	3734	
1	623.0	0.706	4010	
2	747.0	0.879	4727	
3	3270.0	0.940	4973	
4	1544.0	0.524	1454	
...	
2021	11717.0	1.006	2650	
2022	440.0	0.820	2576	
2023	4359.0	0.549	5220	
2024	7359.0	0.558	2094	
2025	20908.0	0.842	8055	

	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio
test_idx			
0	64	0.882	0.504
1	79	0.717	0.716
2	67	0.914	0.569
3	74	0.850	0.317
4	35	0.522	0.584
...
2021	69	0.865	0.125
2022	42	0.750	0.803
2023	75	0.829	0.221
2024	36	0.333	0.155
2025	82	0.673	0.052

[2026 rows x 20 columns]

4 Exploratory Data Analysis

```
[11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8101 entries, 0 to 8100
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CLIENTNUM                            8101 non-null   int64
1   Customer_Age                         8101 non-null   int64
2   Gender                               8101 non-null   object
3   Dependent_count                      8101 non-null   int64
4   Education_Level                      8101 non-null   object
5   Marital_Status                      8101 non-null   object
6   Income_Category                     8101 non-null   object
7   Card_Category                       8101 non-null   object
8   Months_on_book                      8101 non-null   int64
9   Total_Relationship_Count            8101 non-null   int64
10  Months_Inactive_12_mon              8101 non-null   int64
11  Contacts_Count_12_mon               8101 non-null   int64
12  Credit_Limit                        8101 non-null   float64
13  Total_Revolving_Bal                 8101 non-null   int64
14  Avg_Open_To_Buy                     8101 non-null   float64
15  Total_Amt_Chng_Q4_Q1                8101 non-null   float64
16  Total_Trans_Amt                     8101 non-null   int64
17  Total_Trans_Ct                      8101 non-null   int64
18  Total_Ct_Chng_Q4_Q1                 8101 non-null   float64
19  Avg_Utilization_Ratio                8101 non-null   float64
```



```
20 Attrition_Flag          8101 non-null    int64
dtypes: float64(5), int64(11), object(5)
memory usage: 1.4+ MB
```

```
[12]: columns = data.columns
      binary_cols = []
      for col in columns:
          if data[col].value_counts().shape[0] == 2:
              binary_cols.append(col)
      binary_cols
```

```
[12]: ['Gender', 'Attrition_Flag']
```

```
[13]: numeric_cols = data.select_dtypes(include=['number'])
      cat_cols = data.select_dtypes(include=[object])
```

```
[14]: mult_categories_cols = list(set(columns).difference(binary_cols,numeric_cols))
      mult_categories_cols
```

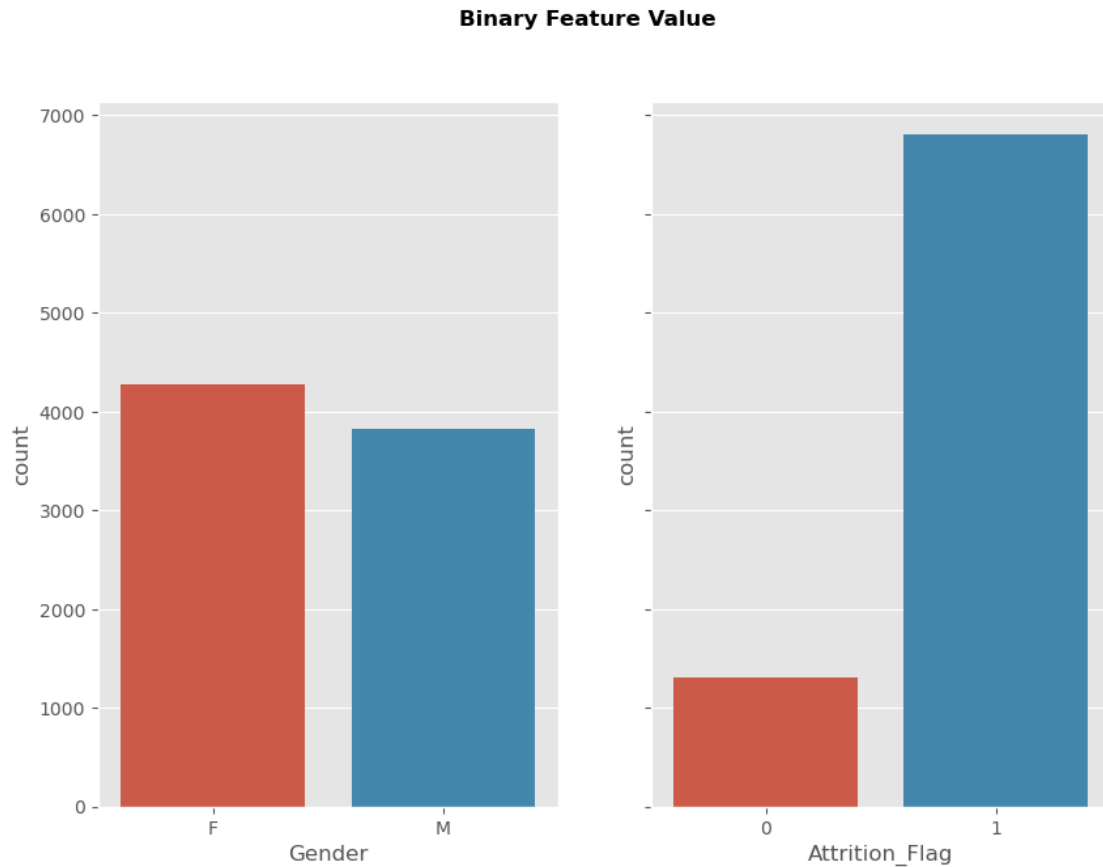
```
[14]: ['Income_Category', 'Card_Category', 'Marital_Status', 'Education_Level']
```

```
[15]: numeric_categories_cols = list(set(numeric_cols).difference(binary_cols))
      numeric_categories_cols
```

```
[15]: ['Total_Trans_Amt',
      'Customer_Age',
      'Avg_Utilization_Ratio',
      'CLIENTNUM',
      'Months_Inactive_12_mon',
      'Total_Amt_Chng_Q4_Q1',
      'Total_Revolving_Bal',
      'Months_on_book',
      'Credit_Limit',
      'Total_Relationship_Count',
      'Total_Trans_Ct',
      'Dependent_count',
      'Contacts_Count_12_mon',
      'Total_Ct_Chng_Q4_Q1',
      'Avg_Open_To_Buy']
```

```
[16]: fig, axes = plt.subplots(1, 2, figsize=(10, 7), sharey=True)

      for i, col in enumerate(binary_cols):
          ax = axes.flat[i]
          sns.countplot(x= data[col],ax=ax)
      plt.suptitle("Binary Feature Value",fontweight='bold')
      plt.show()
```



```
[17]: n_cols = 2

n_rows = plot_template(n_cols,mult_categories_cols)
n_rows
```


```
[17]: 2
```

```
[18]: fig, axes = plt.subplots(n_rows,n_cols, figsize=(3.5 * 4, 3.5 *n_rows),
    ↪sharey=True)

for i, col in enumerate(mult_categories_cols):
    ax = axes.flat[i]
    sns.countplot(x= data[col],ax=ax)

for ax in axes.flatten():
    plt.sca(ax)
    plt.xticks(rotation = 70)
```

```
plt.suptitle("Multi-Feature Values", y=1.05, fontsize=15, fontweight='bold')
plt.tight_layout()
plt.show()
```



```
[19]: n_rows = plot_template(n_cols, numeric_categories_cols)
n_rows
```

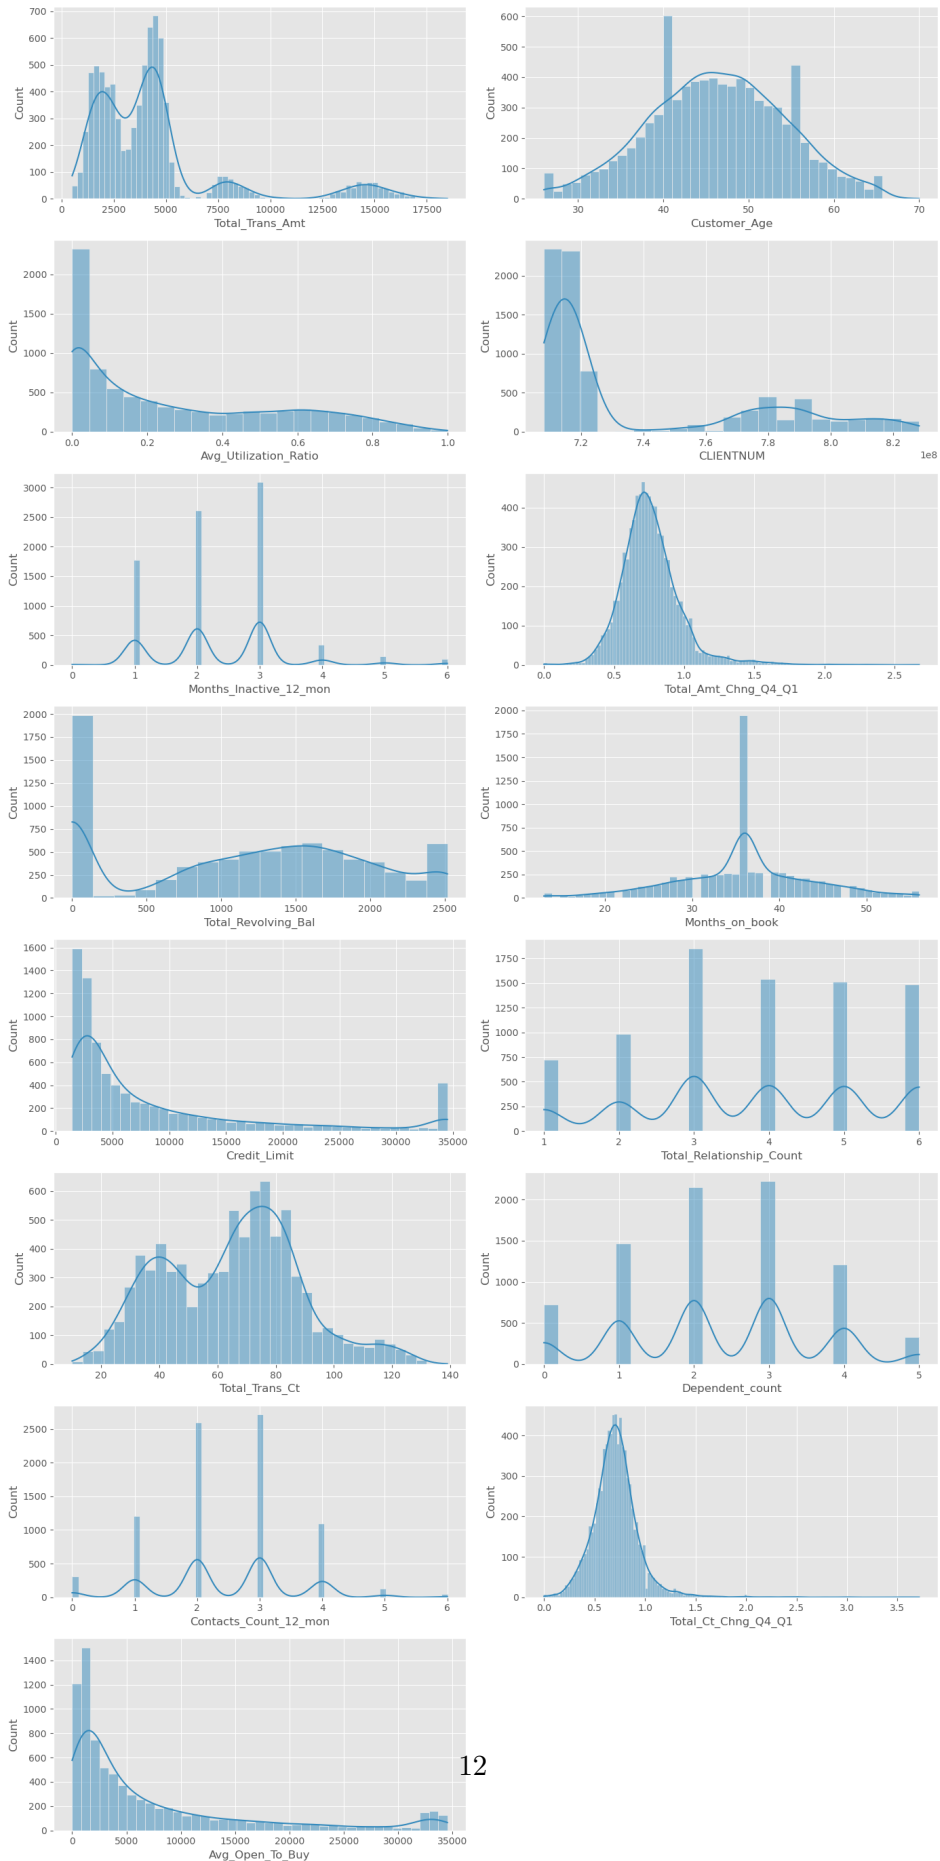
[19]: 8

```
[20]: fig, axes = plt.subplots(n_rows, n_cols, figsize=(3.5 * 4, 3.5 * n_rows))

for i, col in enumerate(numeric_categories_cols):
    ax = axes.flat[i]
    sns.histplot(x= data[col], kde=True, ax=ax)

plt.suptitle("Numeric Values Features", y=1, fontsize=15, fontweight='bold')
axes[-1, -1].axis('off')
plt.tight_layout()
plt.show()
```

Numeric Values Features

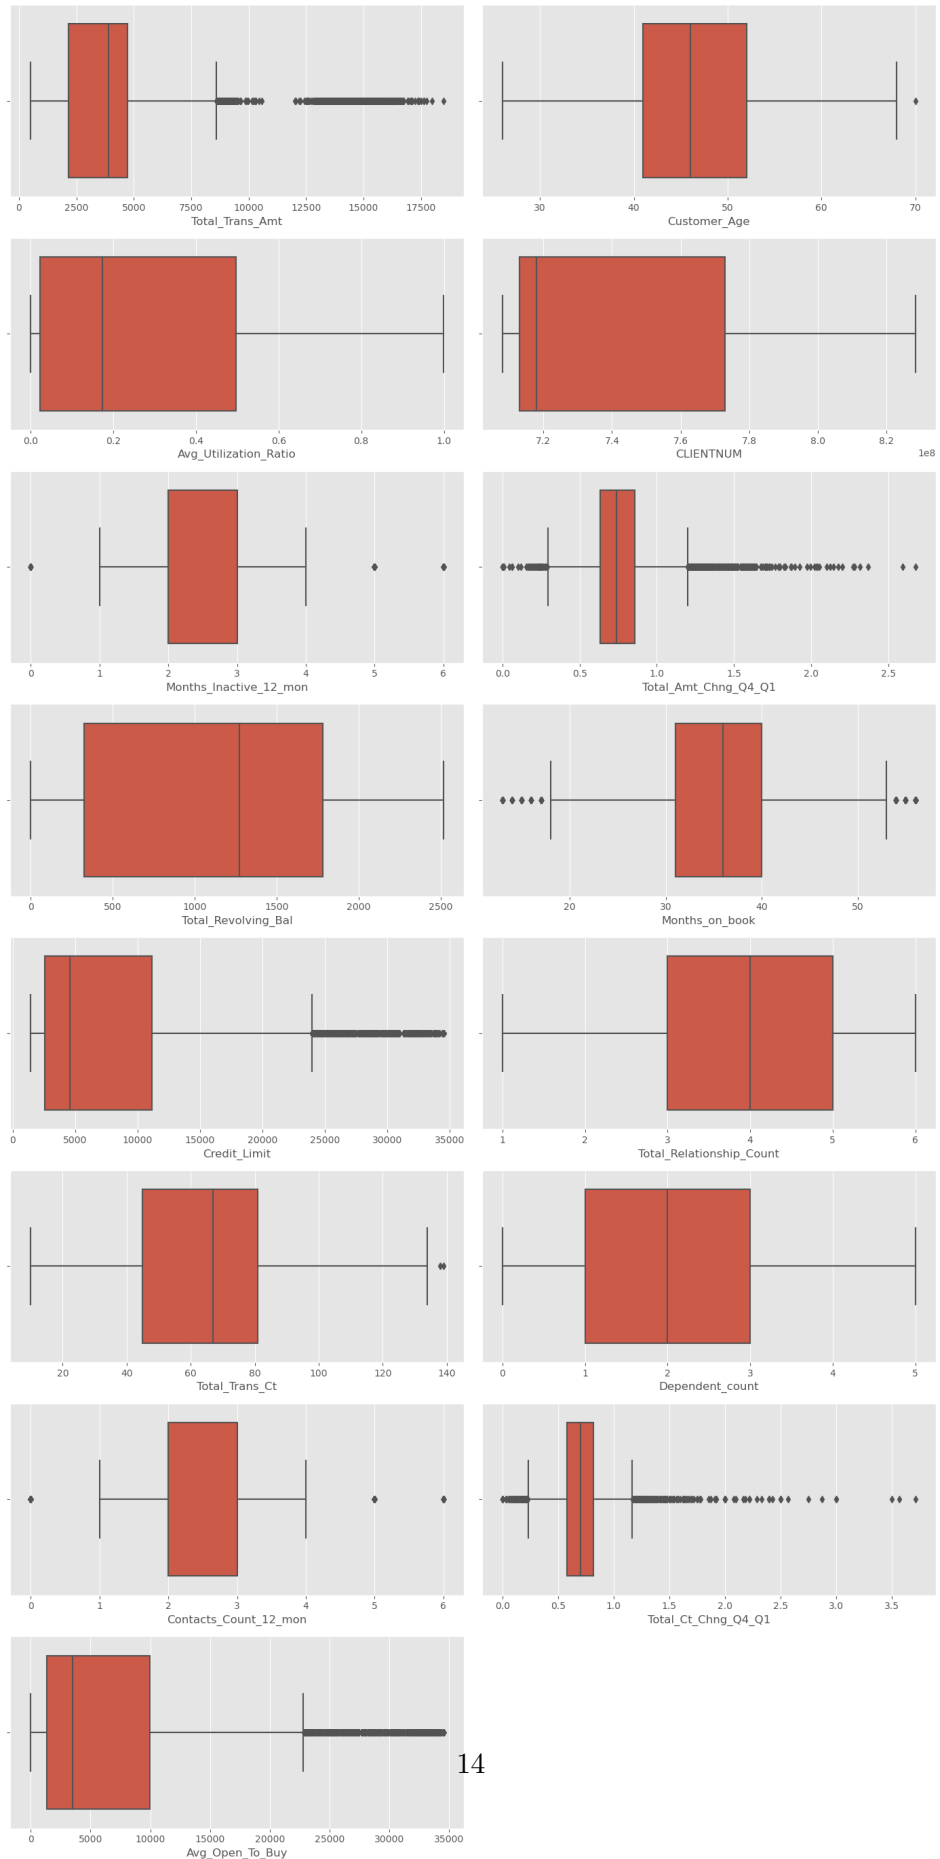


```
[21]: fig, axes = plt.subplots(n_rows,n_cols, figsize=(3.5 * 4, 3.5 *n_rows))

      for i, col in enumerate(numeric_categories_cols):
          ax = axes.flat[i]
          sns.boxplot(x= data[col],ax=ax)

      plt.suptitle("Numeric Values Features", y=1,fontsize=15,fontweight='bold')
      axes[-1, -1].axis('off')
      plt.tight_layout()
      plt.show()
```

Numeric Values Features




It would be interesting to see the behavior of the variables with the attrition flag.

4.1 Categorical Features

```
[22]: gender_attrition = attrited_feature(data, 'Attrition_Flag', 'Gender')
gender_attrition
```

```
[22]:      Attrition_Flag
Gender
F      0.823090
M      0.857928
```

```
[23]: plot_target_attrition(gender_attrition, 'Attrition_Flag')
```

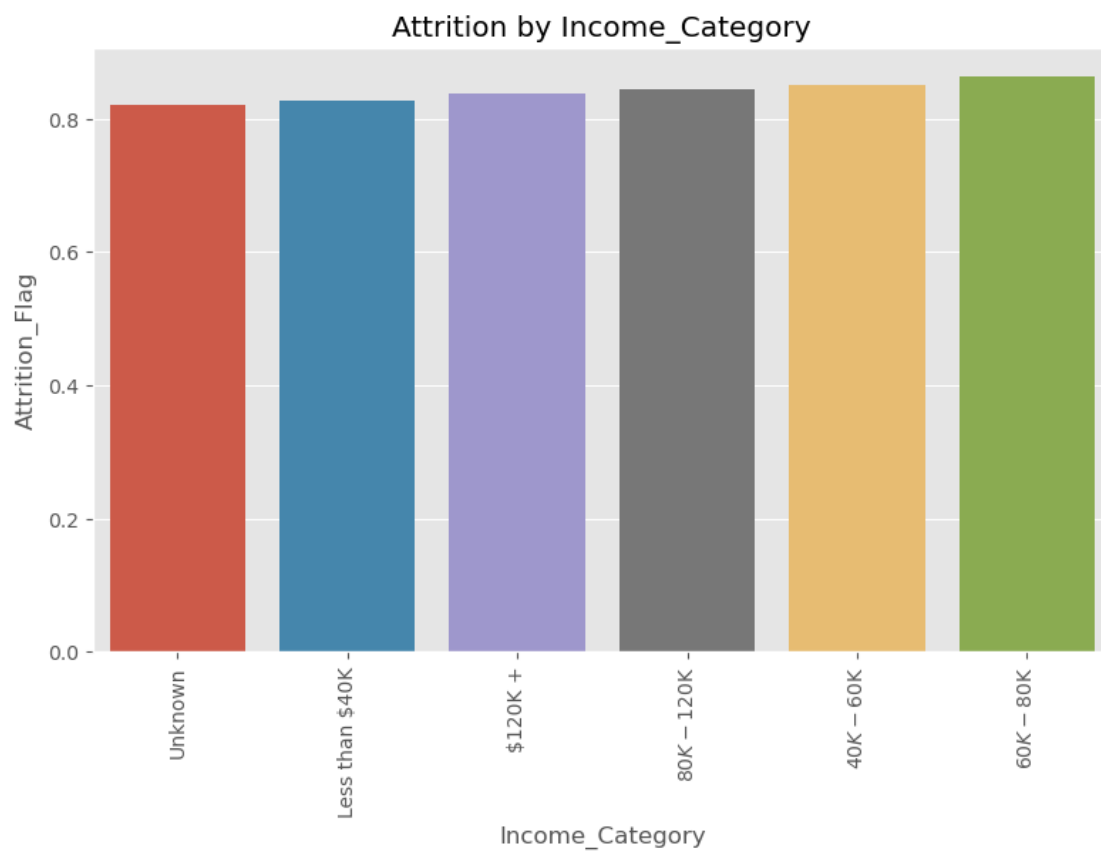


```
[24]: card_category_attrition = \
    attrited_feature(data, 'Attrition_Flag', mult_categories_cols[0])
card_category_attrition
```

```
[24]:      Attrition_Flag
Income_Category
Unknown      0.821147
```

Less than \$40K	0.827881
\$120K +	0.838435
\$80K - \$120K	0.843977
\$40K - \$60K	0.851342
\$60K - \$80K	0.863636

```
[25]: plot_target_attrition(card_category_attrition, 'Attrition_Flag')
```




```
[26]: income_attrition =   
      ↳ attrited_feature(data, 'Attrition_Flag', mult_categories_cols[1])  
      income_attrition
```

```
[26]:
```

Card_Category	Attrition_Flag
Platinum	0.666667
Gold	0.827957
Blue	0.838693
Silver	0.862385


```
[27]: plot_target_attrition(income_attrition, 'Attrition_Flag')
```

```
[28]: education_level_attrition =
      ↳ attrited_feature(data, 'Attrition_Flag', mult_categories_cols[2])
      education_level_attrition
```

```
[28]:
      Attrition_Flag
Marital_Status
Divorced          0.831424
Single            0.833333
Unknown           0.835924
Married           0.846562
```

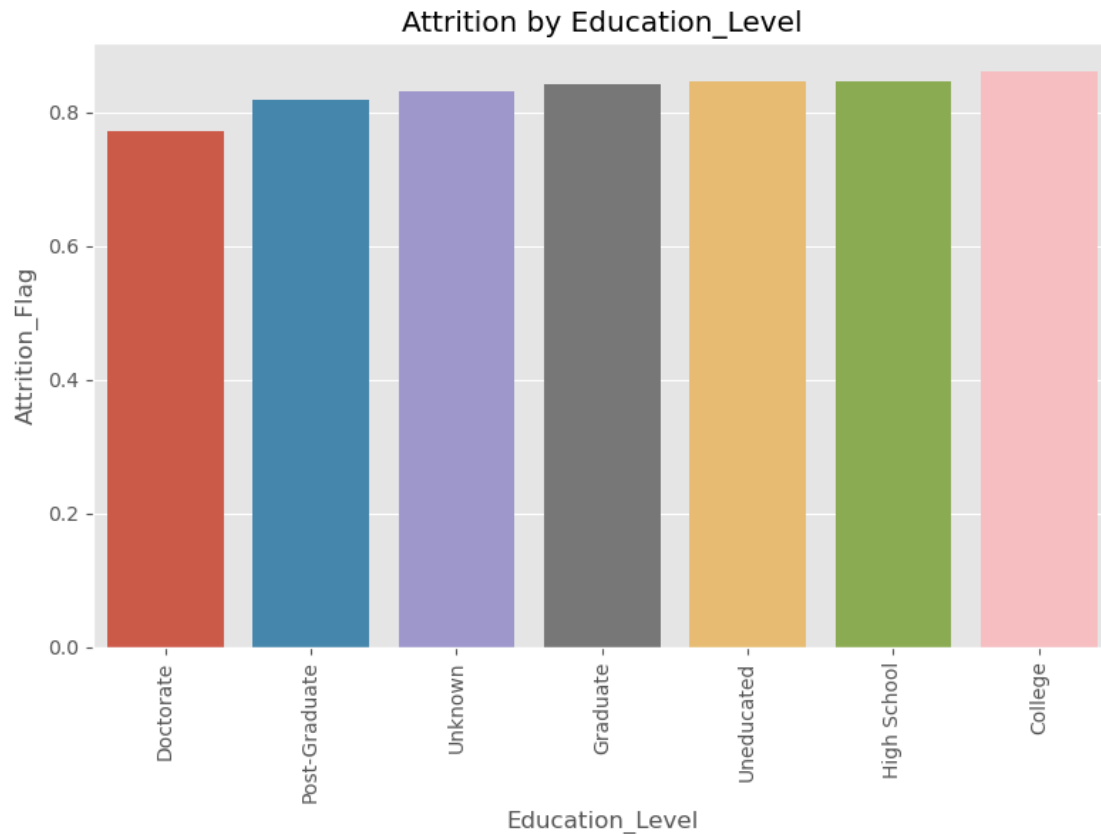
```
[29]: plot_target_attrition(education_level_attrition, 'Attrition_Flag')
```



```
[30]: marital_status_attrition = □
      attrited_feature(data, 'Attrition_Flag', mult_categories_cols[3])
      marital_status_attrition
```

```
[30]: Attrition_Flag
      Education_Level
      Doctorate      0.771831
      Post-Graduate  0.818182
      Unknown        0.831535
      Graduate       0.842168
      Uneducated     0.845431
      High School    0.846819
      College        0.860294
```

```
[31]: plot_target_attrition(marital_status_attrition, 'Attrition_Flag')
```



4.2 Numeric Features

```
[32]: for i in numeric_categories_cols:
      print(i, attrited_feature(data, 'Attrition_Flag', i).describe())
```

Total_Trans_Amt	Attrition_Flag
count	4462.000000
mean	0.823347
std	0.348259
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

Customer_Age	Attrition_Flag
count	44.000000
mean	0.852584
std	0.078781
min	0.500000
25%	0.832231
50%	0.844726

75%	0.874344	
max	1.000000	
Avg_Utilization_Ratio		Attrition_Flag
count	943.000000	
mean	0.881510	
std	0.206084	
min	0.000000	
25%	0.833333	
50%	1.000000	
75%	1.000000	
max	1.000000	
CLIENTNUM		Attrition_Flag
count	8101.000000	
mean	0.839526	
std	0.367068	
min	0.000000	
25%	1.000000	
50%	1.000000	
75%	1.000000	
max	1.000000	
Months_Inactive_12_mon		Attrition_Flag
count	7.000000	
mean	0.755196	
std	0.189499	
min	0.363636	
25%	0.740961	
50%	0.798611	
75%	0.843011	
max	0.956180	
Total_Amt_Chng_Q4_Q1		Attrition_Flag
count	1089.000000	
mean	0.813661	
std	0.266016	
min	0.000000	
25%	0.750000	
50%	0.909091	
75%	1.000000	
max	1.000000	
Total_Revolving_Bal		Attrition_Flag
count	1883.000000	
mean	0.887529	
std	0.260863	
min	0.000000	
25%	1.000000	
50%	1.000000	
75%	1.000000	
max	1.000000	
Months_on_book		Attrition_Flag

count	44.000000	
mean	0.845314	
std	0.044243	
min	0.714286	
25%	0.823432	
50%	0.848812	
75%	0.871630	
max	0.925926	
Credit_Limit		Attrition_Flag
count	5325.000000	
mean	0.845085	
std	0.337039	
min	0.000000	
25%	1.000000	
50%	1.000000	
75%	1.000000	
max	1.000000	
Total_Relationship_Count		Attrition_Flag
count	6.000000	
mean	0.825264	
std	0.074599	
min	0.722843	
25%	0.766316	
50%	0.851069	
75%	0.880784	
max	0.897849	
Total_Trans_Ct		Attrition_Flag
count	126.000000	
mean	0.823723	
std	0.242963	
min	0.000000	
25%	0.724476	
50%	0.955398	
75%	1.000000	
max	1.000000	
Dependent_count		Attrition_Flag
count	6.000000	
mean	0.840738	
std	0.010474	
min	0.827183	
25%	0.835099	
50%	0.839425	
75%	0.845405	
max	0.857338	
Contacts_Count_12_mon		Attrition_Flag
count	7.000000	
mean	0.720160	
std	0.334339	

```

min            0.000000
25%            0.724335
50%            0.793446
75%            0.904311
max            0.990385
Total_Ct_Chng_Q4_Q1    Attrition_Flag
count          795.000000
mean           0.818918
std            0.282779
min            0.000000
25%            0.750000
50%            1.000000
75%            1.000000
max            1.000000
Avg_Open_To_Buy        Attrition_Flag
count          5757.000000
mean           0.842429
std            0.342053
min            0.000000
25%            1.000000
50%            1.000000
75%            1.000000
max            1.000000

```

```

[33]: Dependent_count_attrition =
      ↳ attrited_feature(data, 'Attrition_Flag', 'Dependent_count')
      Dependent_count_attrition

```

```


[33]:                Attrition_Flag
Dependent_count
3                0.827183
4                0.834158
5                0.837920
2                0.840930
0                0.846897
1                0.857338

```

```

[34]: plot_target_attrition(Dependent_count_attrition, 'Attrition_Flag')

```




```
[35]: Contacts_Count_12_mon_attrition = ↳
      attrited_feature(data, 'Attrition_Flag', 'Contacts_Count_12_mon')
      Contacts_Count_12_mon_attrition
```

```
[35]:
```

Contacts_Count_12_mon	Attrition_Flag
6	0.000000
5	0.676692
4	0.771978
3	0.793446
2	0.879045
1	0.929577
0	0.990385

```
[36]: plot_target_attrition(Contacts_Count_12_mon_attrition, 'Attrition_Flag')
```




```
[37]: Total_Trans_Ct_attrition =
      ↳ attrited_feature(data, 'Attrition_Flag', 'Total_Trans_Ct')
      Total_Trans_Ct_attrition
```

```
[37]:
```

Total_Trans_Ct	Attrition_Flag
10	0.000000
12	0.000000
14	0.000000
17	0.166667
15	0.200000
...	...
108	1.000000
109	1.000000
110	1.000000
100	1.000000
139	1.000000

[126 rows x 1 columns]

```
[38]: plot_target_attrition(Total_Trans_Ct_attrition, 'Attrition_Flag')
```





```
[39]: Total_Relationship_Count_attrition = ↳ attrited_feature(data, 'Attrition_Flag', 'Total_Relationship_Count')
Total_Relationship_Count_attrition
```

```
[39]:
```

Total_Relationship_Count	Attrition_Flag
2	0.722843
1	0.746556
3	0.825594
4	0.876543
5	0.882197
6	0.897849

```
[40]: plot_target_attrition(Total_Relationship_Count_attrition, 'Attrition_Flag')
```




```
[41]: Months_on_book_attrition =
↳ attrited_feature(data, 'Attrition_Flag', 'Months_on_book')
Months_on_book_attrition
```

```
[41]: Attrition_Flag
Months_on_book
15          0.714286
50          0.735632
51          0.765625
52          0.788462
39          0.800725
44          0.802139
18          0.804348
30          0.813043
41          0.814815
37          0.818841
22          0.819277
38          0.824818
20          0.825397
48          0.827068
36          0.827179
33          0.832000
56          0.833333
46          0.838710
```

43	0.839623
34	0.842697
45	0.844920
24	0.848000
25	0.849624
28	0.856522
49	0.859649
26	0.861111
42	0.861905
40	0.862454
32	0.863071
47	0.863309
29	0.869792
17	0.870968
35	0.871094
21	0.873239
13	0.877193
55	0.878788
31	0.882353
27	0.883721
54	0.883721
23	0.894737
16	0.900000
53	0.920635
14	0.923077
19	0.925926

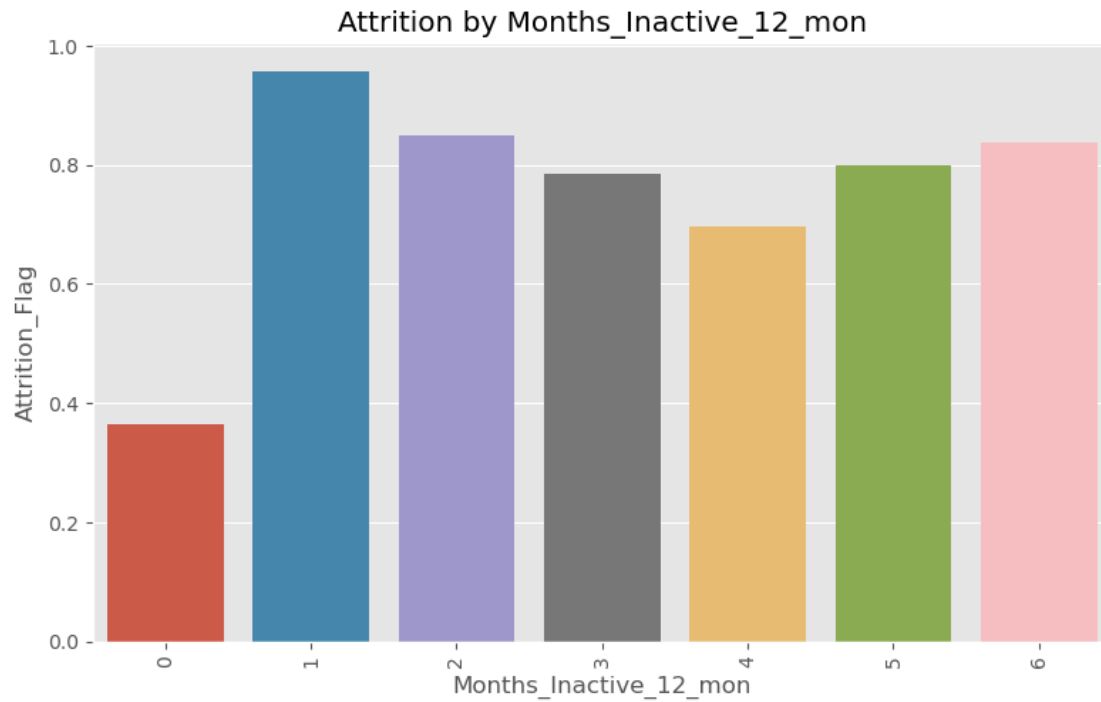
```
[42]: plot_target_attrition(Months_on_book_attrition, 'Attrition_Flag')
```



```
[43]: Months_Inactive_12_mon_attrition =   
      ↪ attrited_feature(data, 'Attrition_Flag', 'Months_Inactive_12_mon')   
      Months_Inactive_12_mon_attrition
```

```
[43]:           Attrition_Flag   
      Months_Inactive_12_mon   
      0                      0.363636   
      4                      0.696532   
      3                      0.785391   
      5                      0.798611   
      6                      0.836538   
      2                      0.849483   
      1                      0.956180
```

```
[44]: plot_target_attrition(Months_Inactive_12_mon_attrition, 'Attrition_Flag')
```




```
[45]: Customer_Age_attrition = attrited_feature(data, 'Attrition_Flag', 'Customer_Age')
      Customer_Age_attrition
```

```
[45]: Attrition_Flag
Customer_Age
68      0.500000
59      0.735537
54      0.779592
61      0.794521
62      0.800000
55      0.811404
41      0.814570
43      0.816712
48      0.817204
30      0.830189
56      0.830189
49      0.832911
47      0.833333
57      0.833333
52      0.833333
44      0.836788
58      0.838462
51      0.840000
46      0.841310
```

40	0.843333
32	0.844444
50	0.844687
39	0.844765
45	0.846154
42	0.846626
38	0.852590
53	0.853821
63	0.862069
27	0.863636
33	0.867347
37	0.872549
34	0.873016
35	0.874126
64	0.875000
31	0.884615
36	0.886905
60	0.888889
65	0.909091
26	0.920635
29	0.940000
28	1.000000
66	1.000000
67	1.000000
70	1.000000

```
[46]: plot_target_attrition(Customer_Age_attrition, 'Attrition_Flag')
```

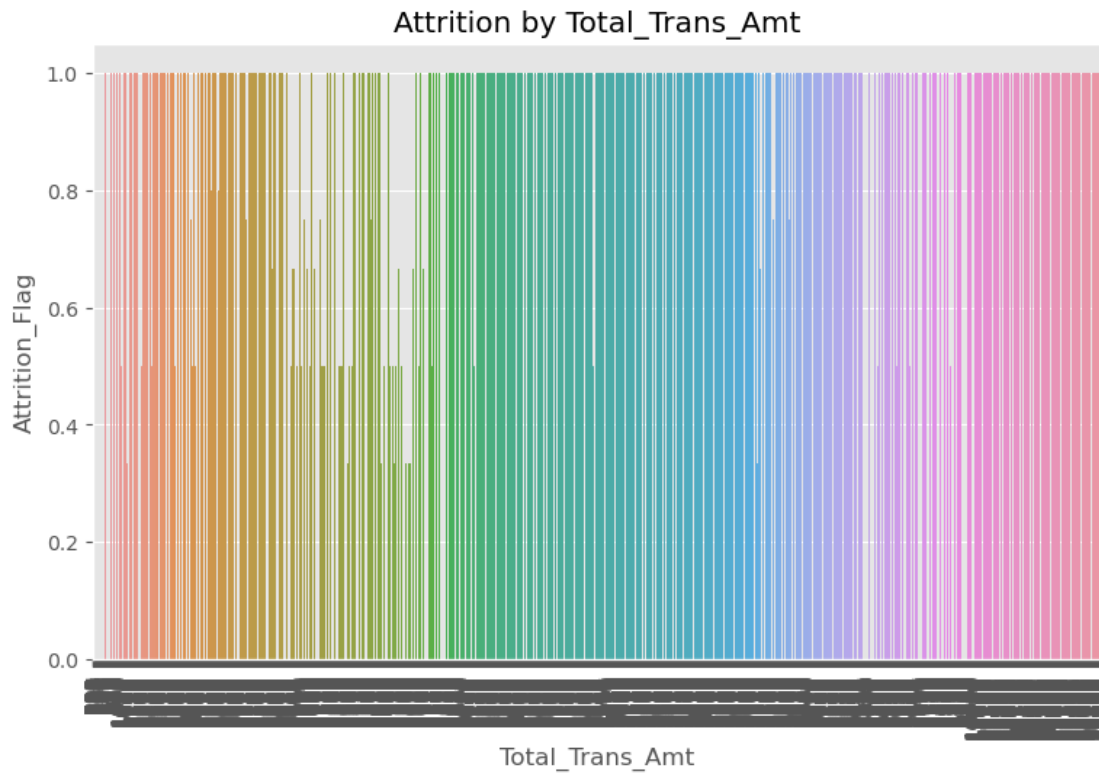


```
[47]: Total_Trans_Amt_attrition =
      attrited_feature(data, 'Attrition_Flag', 'Total_Trans_Amt')
      Total_Trans_Amt_attrition
```

```
[47]: Attrition_Flag
Total_Trans_Amt
510          0.0
2430          0.0
2427          0.0
2426          0.0
2424          0.0
...
3693          1.0
3694          1.0
3695          1.0
3660          1.0
18484         1.0
```

[4462 rows x 1 columns]

```
[48]: plot_target_attrition( Total_Trans_Amt_attrition, 'Attrition_Flag')
```



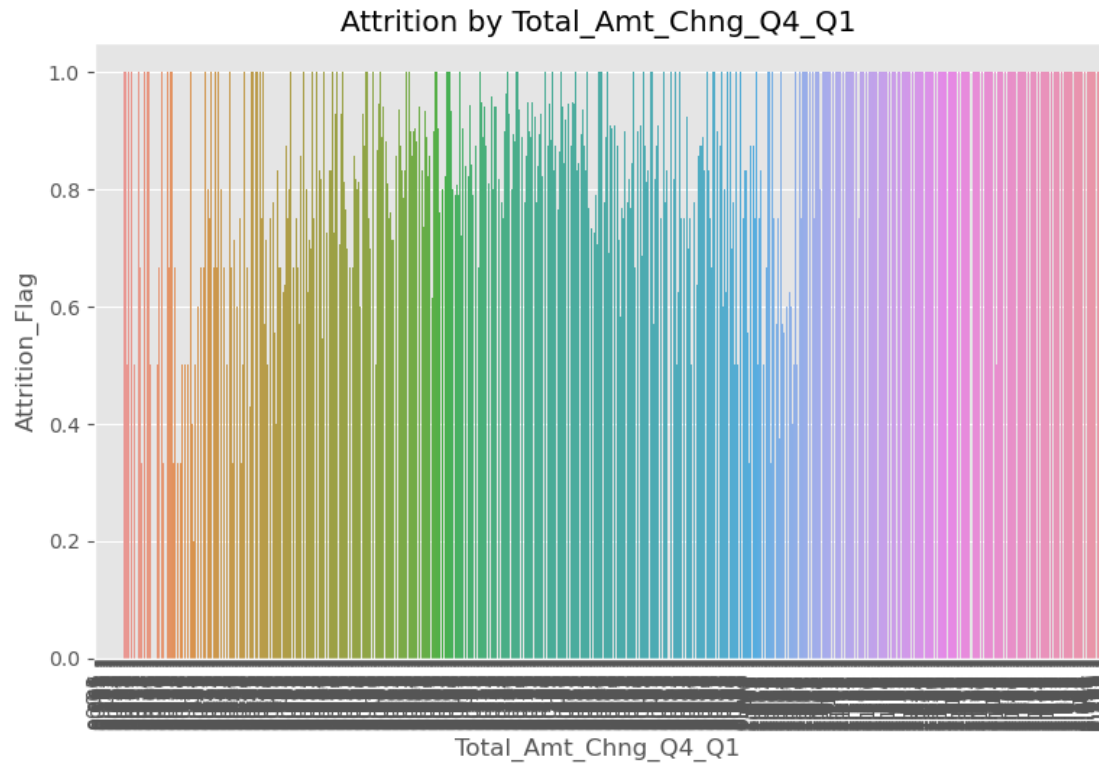
```
[49]: Total_Amt_Chng_Q4_Q1_attrition = ↳ attrited_feature(data, 'Attrition_Flag', 'Total_Amt_Chng_Q4_Q1')
      Total_Amt_Chng_Q4_Q1_attrition
```

```
[49]:
```

Total_Amt_Chng_Q4_Q1	Attrition_Flag
0.000	0.0
0.439	0.0
0.304	0.0
0.305	0.0
1.203	0.0
...	...
1.065	1.0
1.064	1.0
1.063	1.0
1.077	1.0
2.675	1.0

[1089 rows x 1 columns]

```
[50]: plot_target_attrition(Total_Amt_Chng_Q4_Q1_attrition, 'Attrition_Flag')
```

5 Split the Data

```
[51]: X = data.drop(columns=['CLIENTNUM', 'Attrition_Flag'])
      y = data['Attrition_Flag']
      seed = 13
```

```
[52]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=seed)
```

6 Preprocessing Data

```
[53]: cat_cols.columns.to_list()
```

```
[53]: ['Gender',
      'Education_Level',
      'Marital_Status',
      'Income_Category',
      'Card_Category']
```

```
[54]: #encode the categorical values
enc = ce.TargetEncoder(cols=cat_cols.columns,smoothing=11).fit(X_train[cat_cols.
↳ columns],y_train)
X_train[cat_cols.columns] = enc.transform(X_train[cat_cols.columns])
```

```
[55]: X_test[cat_cols.columns] = enc.transform(X_test[cat_cols.columns])
```

```
[56]: X_train.head().T
```

```
[56]: train_idx          802          4121          3083          6160  \
Customer_Age          45.000000          55.000000          29.000000          47.000000
Gender                0.855063          0.824137          0.855063          0.855063
Dependent_count        3.000000          1.000000          0.000000          5.000000
Education_Level        0.843642          0.843642          0.843642          0.859177
Marital_Status         0.828299          0.828299          0.846585          0.846585
Income_Category        0.863014          0.823635          0.863014          0.829218
Card_Category          0.838528          0.838528          0.838528          0.838528
Months_on_book        38.000000          42.000000          36.000000          34.000000
Total_Relationship_Count 1.000000          6.000000          3.000000          3.000000
Months_Inactive_12_mon  3.000000          3.000000          3.000000          1.000000
Contacts_Count_12_mon   1.000000          2.000000          4.000000          2.000000
Credit_Limit         3751.000000         2517.000000         1459.000000         17603.000000
Total_Revolving_Bal    1869.000000         1257.000000          590.000000          1063.000000
Avg_Open_To_Buy       1882.000000         1260.000000          869.000000         16540.000000
Total_Amt_Chng_Q4_Q1    0.926000          0.834000          1.164000          1.065000
Total_Trans_Amt       9497.000000         4528.000000         2229.000000          2902.000000
Total_Trans_Ct         106.000000          75.000000          52.000000          64.000000
Total_Ct_Chng_Q4_Q1    0.710000          0.786000          1.080000          0.641000
Avg_Utilization_Ratio   0.498000          0.499000          0.404000          0.060000

train_idx          509
Customer_Age          41.000000
Gender                0.855063
Dependent_count        2.000000
Education_Level        0.836798
Marital_Status         0.828299
Income_Category        0.838614
Card_Category          0.838528
Months_on_book        28.000000
Total_Relationship_Count 1.000000
Months_Inactive_12_mon  2.000000
Contacts_Count_12_mon   3.000000
Credit_Limit         13679.000000
Total_Revolving_Bal    2267.000000
Avg_Open_To_Buy       11412.000000
Total_Amt_Chng_Q4_Q1    1.232000
Total_Trans_Amt       15180.000000
```

Total_Trans_Ct	103.000000
Total_Ct_Chng_Q4_Q1	0.689000
Avg_Utilization_Ratio	0.166000

```
[57]: #scaler = StandardScaler()
      scaler = RobustScaler()

      # transform data
      X_train[X_train.columns] = scaler.fit_transform(X_train)
      X_test[X_test.columns] = scaler.transform(X_test)
```

```
[58]: n_rows = plot_template(n_cols,X_train.columns)
      n_rows
```

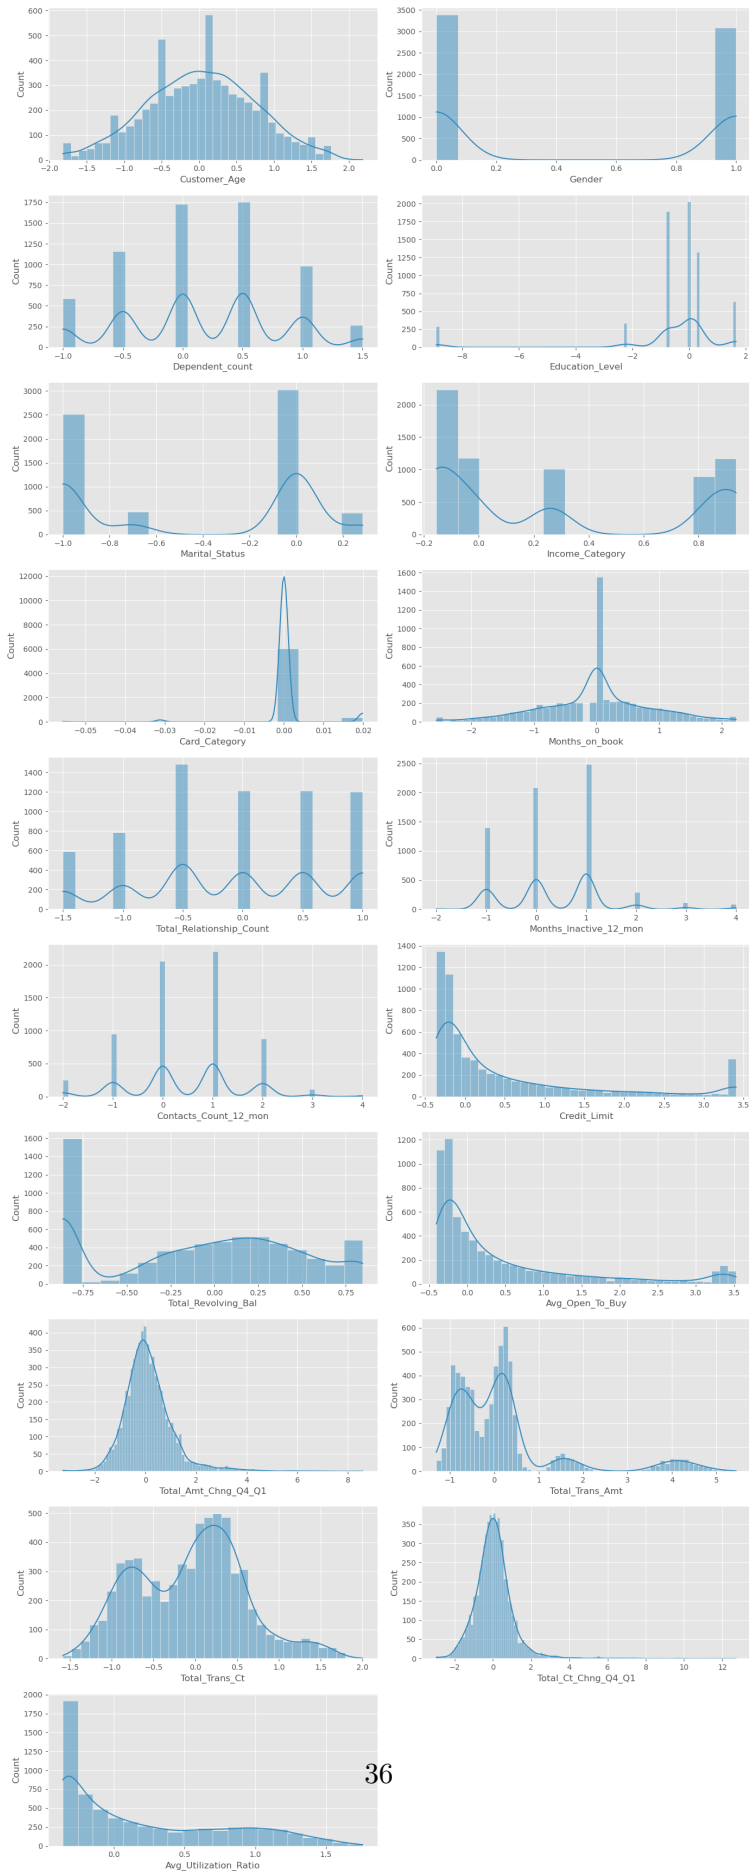
```
[58]: 10
```

```
[59]: fig, axes = plt.subplots(n_rows,n_cols, figsize=(3.5 * 4, 3.5 *n_rows))

      for i, col in enumerate(X_train):
          ax = axes.flat[i]
          sns.histplot(x= X_train[col],kde=True,ax=ax)

      plt.suptitle("Numeric Values Features", y=1,fontsize=15,fontweight='bold')
      axes[-1, -1].axis('off')
      plt.tight_layout()
      plt.show()
```

Numeric Values Features

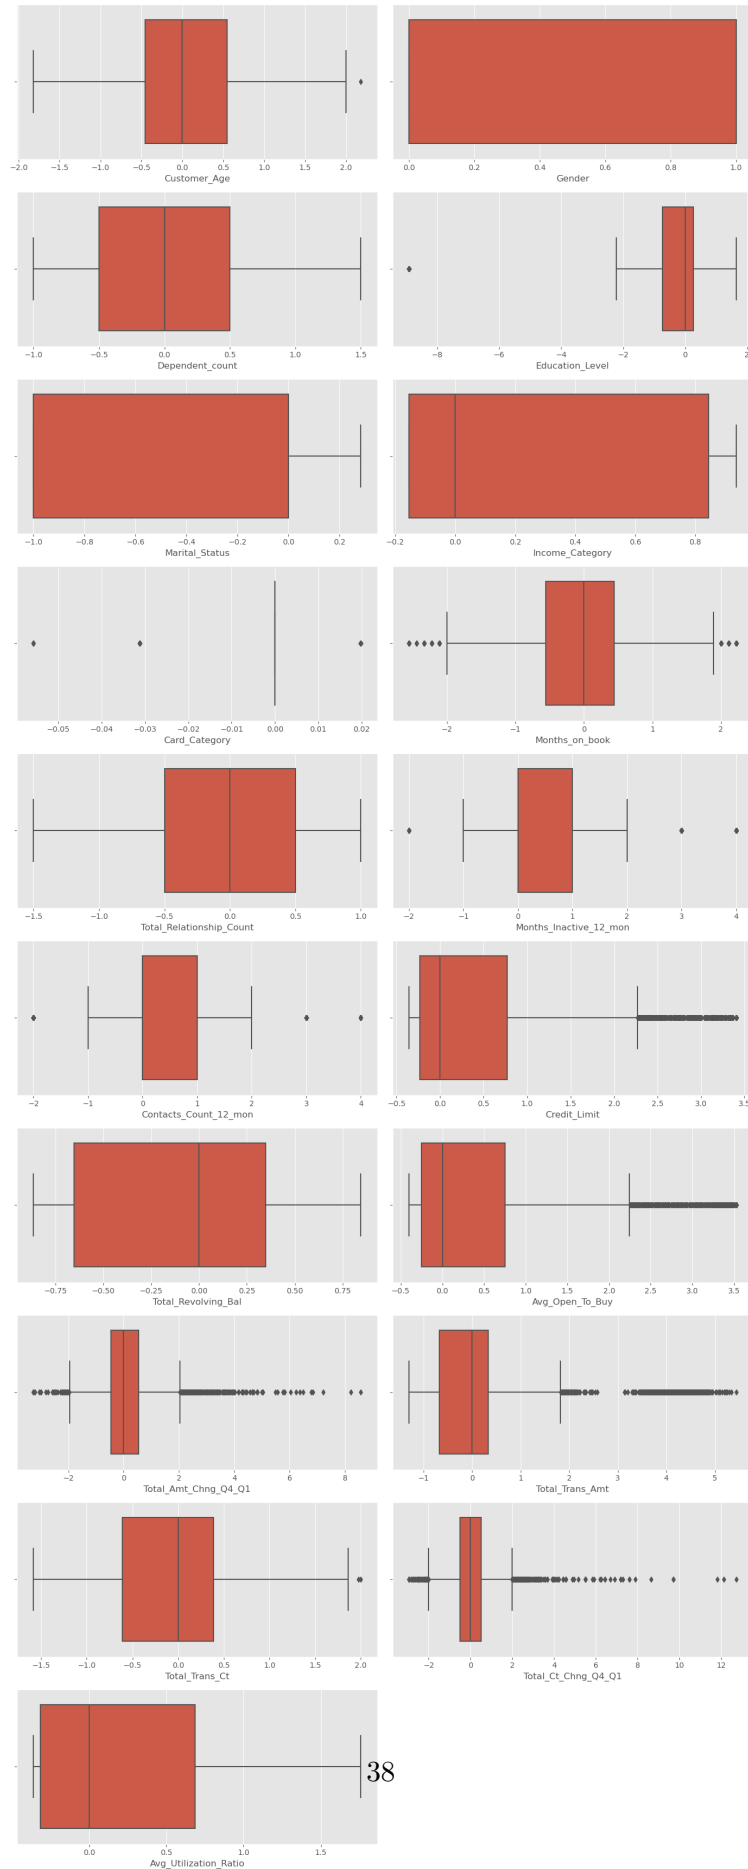


```
[60]: fig, axes = plt.subplots(n_rows,n_cols, figsize=(3.5 * 4, 3.5 *n_rows))

for i, col in enumerate(X_train.columns):
    ax = axes.flat[i]
    sns.boxplot(x= X_train[col],ax=ax)

plt.suptitle("Numeric Values Features", y=1,fontsize=15,fontweight='bold')
axes[-1, -1].axis('off')
plt.tight_layout()
plt.show()
```

Numeric Values Features



```
# Balance Data
```

```
[61]: over = SMOTE(sampling_strategy='not majority',random_state=17)
      under = RandomUnderSampler(sampling_strategy='not majority')
      steps = [ ('o', over), ('u', under)]
      pipeline = Pipeline(steps=steps)

      X_train_sm,y_train_sm = pipeline.fit_resample(X_train,y_train)
```

```
[62]: X_train_sm.shape[0],y_train_sm.shape[0],X_test.shape[0],y_test.shape[0]
```


```
[62]: (10872, 10872, 1621, 1621)
```

7 Feature Selection

```
[63]: features = XGBClassifier()
      features.fit(X_train_sm, y_train_sm)

      plt.figure(figsize=(10,6))
      plot_importance(features,height=0.3,importance_type='gain')
      plt.tight_layout()
      plt.show()
```

<Figure size 1000x600 with 0 Axes>



```
[64]: features_importance_names = features.get_booster().
      ↪get_score(importance_type='gain')
      features_importance_names
```

```
[64]: {'Customer_Age': 5.369898319244385,
      'Gender': 12.305244445800781,
      'Dependent_count': 2.7322099208831787,
      'Education_Level': 2.398404359817505,
      'Marital_Status': 4.644156455993652,
      'Income_Category': 2.046273708343506,
      'Card_Category': 0.8429279923439026,
      'Months_on_book': 1.6233980655670166,
      'Total_Relationship_Count': 15.567176818847656,
      'Months_Inactive_12_mon': 7.974614143371582,
      'Contacts_Count_12_mon': 6.786797046661377,
      'Credit_Limit': 2.399428129196167,
      'Total_Revolving_Bal': 15.522917747497559,
      'Avg_Open_To_Buy': 1.8516833782196045,
      'Total_Amt_Chng_Q4_Q1': 3.2221853733062744,
      'Total_Trans_Amt': 10.984789848327637,
      'Total_Trans_Ct': 51.22013854980469,
```



```
'Total_Ct_Chng_Q4_Q1': 7.49886417388916,  
'Avg_Utilization_Ratio': 1.1399104595184326}
```

```
[65]: #use function (5) to get the tuple from the list and then sort it  
features_importance_names_list = list(features_importance_names.items())  
features_importance_names_list.sort(key = get_value,reverse=True)
```

```
[66]: features_importance_names_list
```

```
[66]: [('Total_Trans_Ct', 51.22013854980469),  
      ('Total_Relationship_Count', 15.567176818847656),  
      ('Total_Revolving_Bal', 15.522917747497559),  
      ('Gender', 12.305244445800781),  
      ('Total_Trans_Amt', 10.984789848327637),  
      ('Months_Inactive_12_mon', 7.974614143371582),  
      ('Total_Ct_Chng_Q4_Q1', 7.49886417388916),  
      ('Contacts_Count_12_mon', 6.786797046661377),  
      ('Customer_Age', 5.369898319244385),  
      ('Marital_Status', 4.644156455993652),  
      ('Total_Amt_Chng_Q4_Q1', 3.2221853733062744),  
      ('Dependent_count', 2.7322099208831787),  
      ('Credit_Limit', 2.399428129196167),  
      ('Education_Level', 2.398404359817505),  
      ('Income_Category', 2.046273708343506),  
      ('Avg_Open_To_Buy', 1.8516833782196045),  
      ('Months_on_book', 1.6233980655670166),  
      ('Avg_Utilization_Ratio', 1.1399104595184326),  
      ('Card_Category', 0.8429279923439026)]
```

```
[67]: #get only the first value of the tuple  
features_sorted_names = []  
  
for i in features_importance_names_list:  
    features_sorted_names.append(i[0])  
  
features_sorted_names
```

```
[67]: ['Total_Trans_Ct',  
      'Total_Relationship_Count',  
      'Total_Revolving_Bal',  
      'Gender',  
      'Total_Trans_Amt',  
      'Months_Inactive_12_mon',  
      'Total_Ct_Chng_Q4_Q1',  
      'Contacts_Count_12_mon',  
      'Customer_Age',  
      'Marital_Status',
```

```

'Total_Amt_Chng_Q4_Q1',
'Dependent_count',
'Credit_Limit',
'Education_Level',
'Income_Category',
'Avg_Open_To_Buy',
'Months_on_book',
'Avg_Utilization_Ratio',
'Card_Category']

```

```

[74]: #define the columns with least value
drop_cols = features_sorted_names[9:]
drop_cols

```

```

[74]: ['Marital_Status',
'Total_Amt_Chng_Q4_Q1',
'Dependent_count',
'Credit_Limit',
'Education_Level',
'Income_Category',
'Avg_Open_To_Buy',
'Months_on_book',
'Avg_Utilization_Ratio',
'Card_Category']

```

```

[75]: X_train_ = X_train_sm.drop(columns=drop_cols)
y_train_ = y_train_sm

```

```

[76]: X_test_ = X_test.drop(columns=drop_cols)
y_test_ = y_test

```

8 Model Creation

8.1 Random Forest

```

[77]: rf = RandomForestClassifier(random_state= seed)

```

```

[78]: # Hyperparameter Tuning
parameters_rf = {
    'n_estimators': [10,100, 200,500],
    'max_depth' : [2, 5, 7,10,20,25],
    'min_samples_split': [12,20,40],
    'criterion':['gini', 'entropy'],
    'bootstrap': [True, False]
}

```

```
[79]: skfold = StratifiedKFold(n_splits=6,shuffle=True,random_state=seed)
```

```
[80]: gridRf = GridSearchCV(rf, parameters_rf, cv=skfold)
```

```
[81]: y_rf_score = gridRf.fit(X_train_, y_train_)
```

```
[82]: gridRf.best_score_
```

```
[82]: 0.9734179543782192
```

```
[83]: #best parameters
      gridRf.best_params_
```

```
[83]: {'bootstrap': False,
      'criterion': 'entropy',
      'max_depth': 20,
      'min_samples_split': 12,
      'n_estimators': 500}
```

```
[84]: best_estimator_grid_rf = gridRf.best_estimator_
      best_estimator_grid_rf
```

```
[84]: RandomForestClassifier(bootstrap=False, criterion='entropy', max_depth=20,
                             min_samples_split=12, n_estimators=500, random_state=13)
```

```
[85]: y_rf_pred = gridRf.predict(X_test_)
```

```
[86]: report_rf = classification_report(y_test_, y_rf_pred,output_dict=True)
      df_report_rf = pd.DataFrame(report_rf).T
      df_report_rf
```

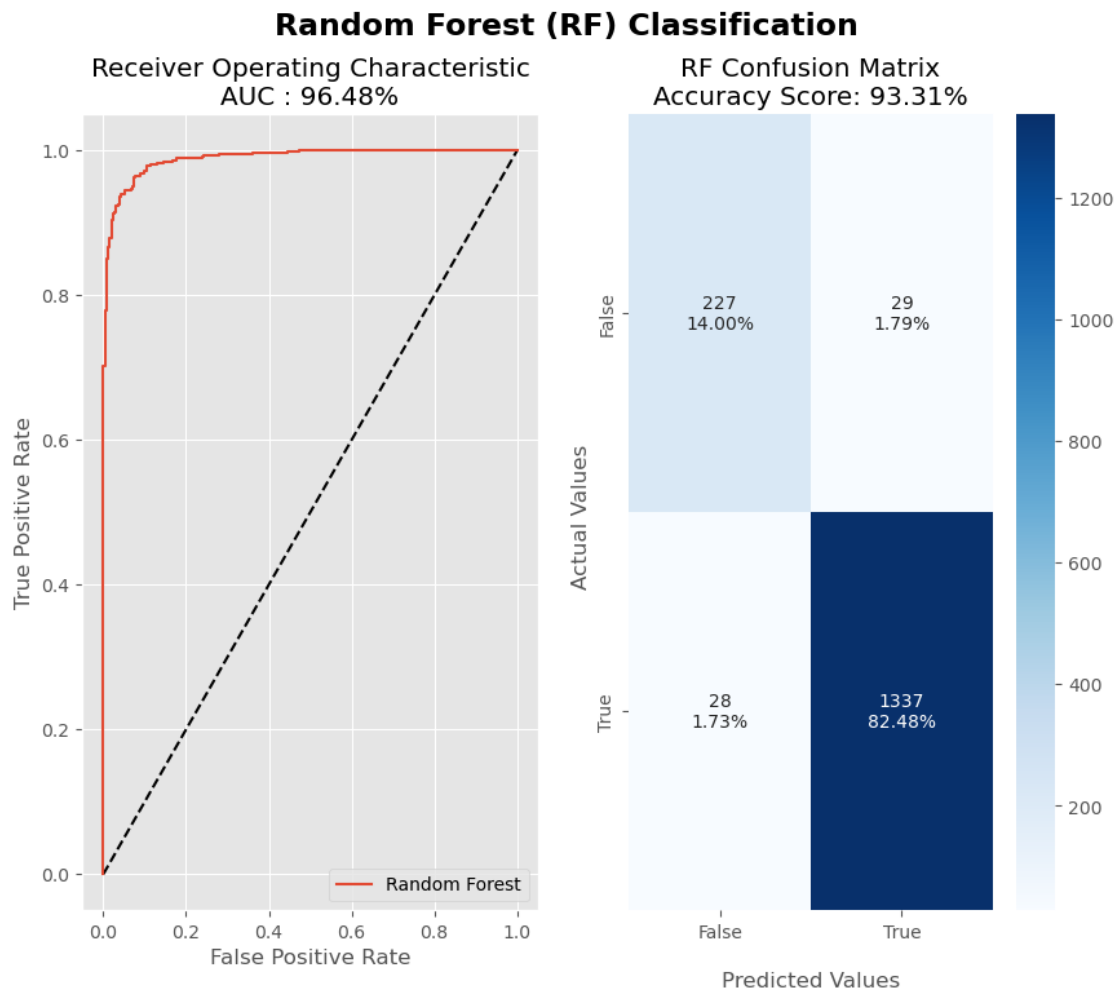
```
[86]:
```

	precision	recall	f1-score	support
0	0.890196	0.886719	0.888454	256.000000
1	0.978770	0.979487	0.979129	1365.000000
accuracy	0.964837	0.964837	0.964837	0.964837
macro avg	0.934483	0.933103	0.933791	1621.000000
weighted avg	0.964782	0.964837	0.964809	1621.000000

```
[87]: #Confusion matrix
      conf_matrix_rf = confusion_matrix(y_test_, y_rf_pred)
      #Accuracy
      acu_rf = accuracy_score(y_test_, y_rf_pred)
      #AUC
      auc_rf = roc_auc_score(y_test_, y_rf_pred)
```

```
[88]: acmod_rf = {"RF_y_test":round(acu_rf,4)}
      aucmod_rf = {"RF_y_test":round(auc_rf,4)}
```

```
[89]: #run the functiona (4)
subplots_ROC_CM("Random Forest (RF) Classification","Receiver Operating_
↳Characteristic",
               "RF Confusion Matrix",
               gridRf.
↳best_estimator_,acmod_rf,aucmod_rf,conf_matrix_rf,"RF_y_test",X_test_,y_test_, "Random_
↳Forest")
```



8.2 Stochastic Gradient Descent (SDG)

```
[90]: sdg = SGDClassifier(random_state= seed)
```

```
[91]: # Hyperparameter Tuning
parameters_sdg = {
    'penalty':['l2', 'l1', 'elasticnet'],
    'learning_rate': ['constant','optimal','invscaling','adaptive'],
```

```

    'loss': ['log_loss', 'modified_huber'],
    'alpha': list(range(1,50,1)) ,
    'fit_intercept': [True, False],
    'max_iter': [5000]

}

```

```
[92]: gridSdg = GridSearchCV(sdg, parameters_sdg, cv=skfold)
```

```
[93]: y_sdg_score = gridSdg.fit(X_train_, y_train_)
```

```

/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:
12348 fits failed out of a total of 14112.
The score on these train-test partitions for these parameters will be set to
nan.
If these failures are not expected, you can try to debug them by setting
error_score='raise'.

```

Below are more details about the failures:

```

-----
10584 fits failed with the following error:
Traceback (most recent call last):
  File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 680, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/linear_model/_stochastic_gradient.py", line 883, in fit
    return self._fit(
  File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/linear_model/_stochastic_gradient.py", line 649, in _fit
    self._validate_params()
  File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/linear_model/_stochastic_gradient.py", line 149, in
_validate_params
    raise ValueError("eta0 must be > 0")
ValueError: eta0 must be > 0

```

```

-----
1764 fits failed with the following error:
Traceback (most recent call last):
  File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 680, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/linear_model/_stochastic_gradient.py", line 883, in fit
    return self._fit(

```

```

File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/linear_model/_stochastic_gradient.py", line 649, in _fit
    self._validate_params()
File "/opt/anaconda3/lib/python3.9/site-
packages/sklearn/linear_model/_stochastic_gradient.py", line 162, in
_validate_params
    raise ValueError("The loss %s is not supported. " % self.loss)
ValueError: The loss log_loss is not supported.

warnings.warn(some_fits_failed_message, FitFailedWarning)
/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_search.py:969: UserWarning: One or more of the
test scores are non-finite: [nan nan nan ... nan nan nan]
    warnings.warn(

```

```
[94]: gridSdg.best_score_
```

```
[94]: 0.8084988962472406
```

```
[95]: #best parameters
      gridSdg.best_params_
```

```
[95]: {'alpha': 1,
      'fit_intercept': True,
      'learning_rate': 'optimal',
      'loss': 'modified_huber',
      'max_iter': 5000,
      'penalty': 'l2'}
```

```
[96]: best_estimator_grid_sdg = gridSdg.best_estimator_
      best_estimator_grid_sdg
```

```
[96]: SGDClassifier(alpha=1, loss='modified_huber', max_iter=5000, random_state=13)
```

```
[97]: y_sdg_pred = gridSdg.predict(X_test_)
```

```
[98]: report_sdg = classification_report(y_test_, y_sdg_pred,output_dict=True)
      df_report_sdg = pd.DataFrame(report_sdg).T
      df_report_sdg
```

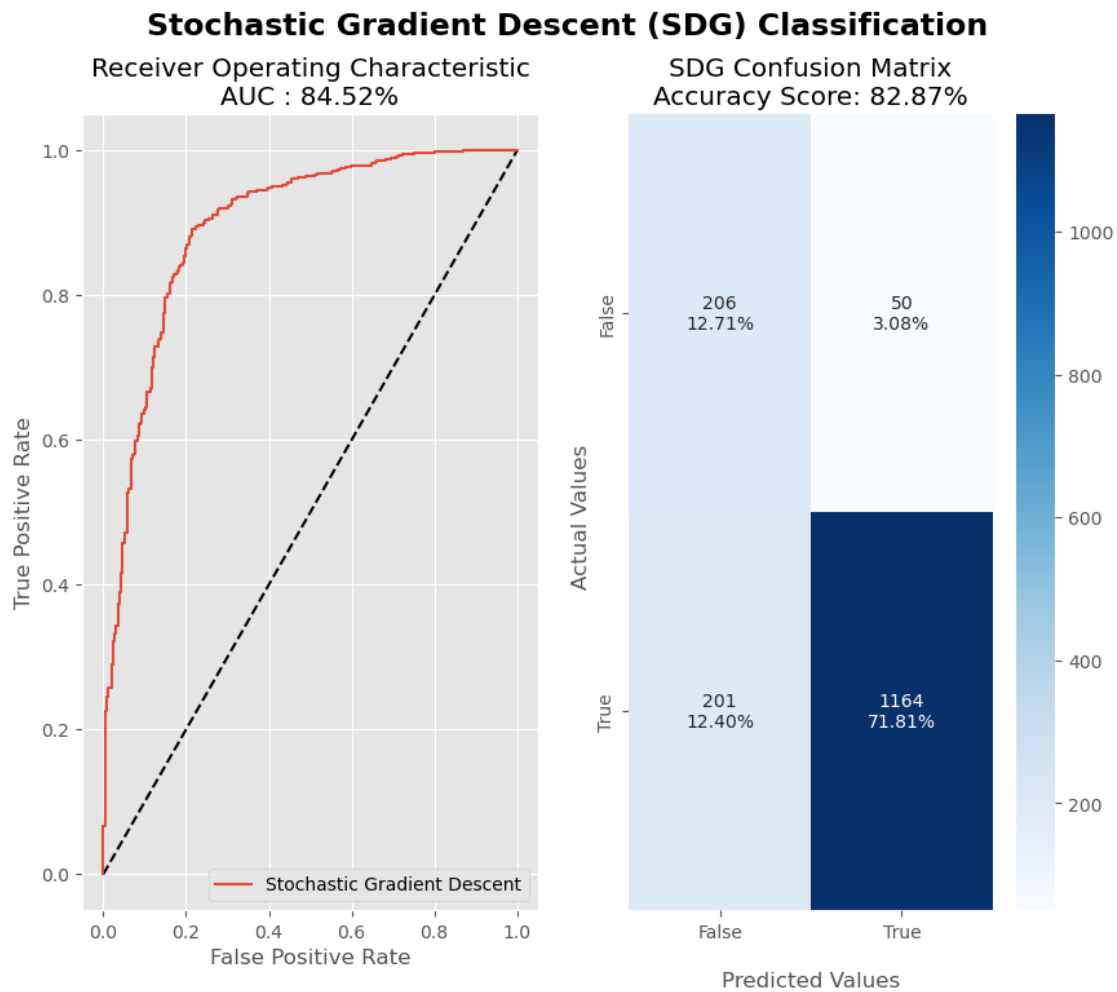
```
[98]:
```

	precision	recall	f1-score	support
0	0.506143	0.804688	0.621418	256.000000
1	0.958814	0.852747	0.902675	1365.000000
accuracy	0.845157	0.845157	0.845157	0.845157
macro avg	0.732478	0.828717	0.762047	1621.000000
weighted avg	0.887325	0.845157	0.858257	1621.000000

```
[99]: #Confusion matrix
conf_matrix_sdg = confusion_matrix(y_test, y_sdg_pred)
#Accuracy
acu_sdg = accuracy_score(y_test, y_sdg_pred)
#AUC
auc_sdg = roc_auc_score(y_test, y_sdg_pred)
```

```
[100]: acmod_sdg = {"SDG_y_test":round(acu_sdg,4)}
aucmod_sdg = {"SDG_y_test":round(auc_sdg,4)}
```

```
[101]: #run the functiona (4)
subplots_ROC_CM("Stochastic Gradient Descent (SDG) Classification","Receiver_
↪Operating Characteristic",
                "SDG Confusion Matrix",
                gridSdg.
↪best_estimator_,acmod_sdg,aucmod_sdg,conf_matrix_sdg,"SDG_y_test",X_test_,y_test_,"Stochast
↪Gradient Descent")
```



8.3 XGBoost

```
[102]: xgbc = XGBClassifier(random_state= seed)
```

```
[103]: # Hyperparameter Tuning
parameters_xgbc = {
    'objective': ['binary:logistic'],
    'learning_rate': [0.0001, 0.001, 0.01, 0.1],
    'max_depth': [2, 4, 7, 15],
    'subsample': [0.1, 0.5, 0.8],
    'colsample_bytree': [0.1, 0.5, 0.7],
    'n_estimators': [100, 300, 500, 1000],
}
```

```
[104]: gridXgbc = GridSearchCV(xgbc, parameters_xgbc, cv=skfold)
```

```
[105]: y_xgbc_score = gridXgbc.fit(X_train_, y_train_)
```

```
[106]: gridXgbc.best_score_
```

```
[106]: 0.978476821192053
```

```
[107]: #best parameters
gridXgbc.best_params_
```

```
[107]: {'colsample_bytree': 0.5,
      'learning_rate': 0.1,
      'max_depth': 15,
      'n_estimators': 500,
      'objective': 'binary:logistic',
      'subsample': 0.8}
```

```
[108]: best_estimator_grid_xgbc = gridXgbc.best_estimator_
best_estimator_grid_xgbc
```

```
[108]: XGBClassifier(base_score=None, booster=None, callbacks=None,
      colsample_bylevel=None, colsample_bynode=None,
      colsample_bytree=0.5, early_stopping_rounds=None,
      enable_categorical=False, eval_metric=None, feature_types=None,
      gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
      interaction_constraints=None, learning_rate=0.1, max_bin=None,
      max_cat_threshold=None, max_cat_to_onehot=None,
      max_delta_step=None, max_depth=15, max_leaves=None,
      min_child_weight=None, missing=nan, monotone_constraints=None,
      n_estimators=500, n_jobs=None, num_parallel_tree=None,
```



```
predictor=None, random_state=13, ...)
```

```
[109]: y_xgbc_pred = gridXgbc.predict(X_test_)
```

```
[110]: report_xgbc = classification_report(y_test_, y_xgbc_pred,output_dict=True)
df_report_xgbc = pd.DataFrame(report_xgbc).T
df_report_xgbc
```

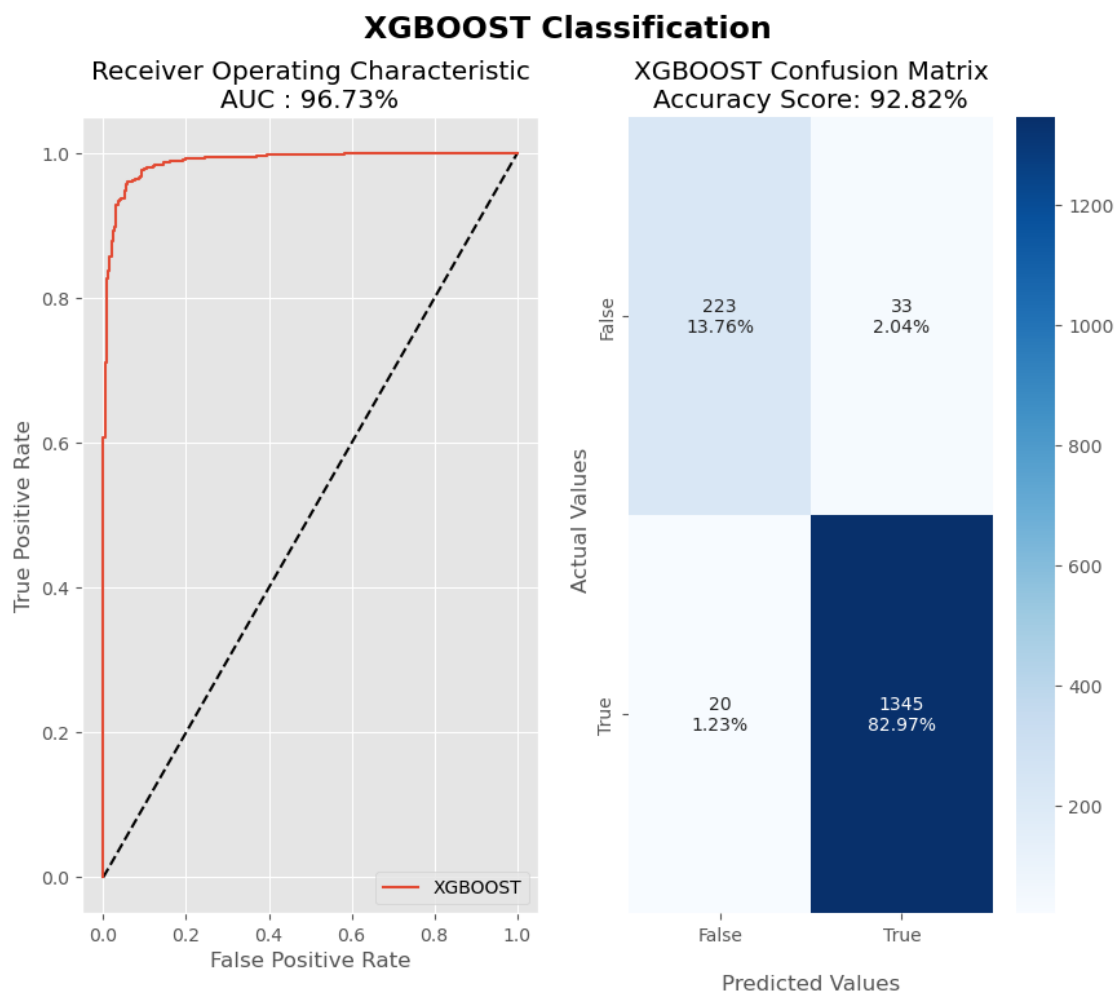
```
[110]:
```

	precision	recall	f1-score	support
0	0.917695	0.871094	0.893788	256.000000
1	0.976052	0.985348	0.980678	1365.000000
accuracy	0.967304	0.967304	0.967304	0.967304
macro avg	0.946874	0.928221	0.937233	1621.000000
weighted avg	0.966836	0.967304	0.966956	1621.000000

```
[111]: #Confusion matrix
conf_matrix_xgbc = confusion_matrix(y_test_, y_xgbc_pred)
#Accuracy
acu_xgbc = accuracy_score(y_test_, y_xgbc_pred)
#AUC
auc_xgbc = roc_auc_score(y_test_, y_xgbc_pred)
```

```
[112]: acmod_xgbc = {"XGBC_y_test":round(acu_xgbc,4)}
aucmod_xgbc = {"XGBC_y_test":round(auc_xgbc,4)}
```

```
[113]: #run the functiona (9)
subplots_ROC_CM("XGBOOST Classification","Receiver Operating Characteristic",
                "XGBOOST Confusion Matrix",
                gridXgbc.
                ↪best_estimator_,acmod_xgbc,aucmod_xgbc,conf_matrix_xgbc,"XGBC_y_test",X_test_,y_test_,"XGB0
```



9 Select Model

```
[138]: Accuracy = {"RF":acu_rf,"SDG":acu_sdg,"XGBC":acu_xgbc}
AUC = {"RF":auc_rf,"SDG":auc_sdg,"XGBC":auc_xgbc}
F1_score = {"RF":df_report_rf.iloc[3,2],
            "SDG":df_report_sdg.iloc[3,2],"XGBC":df_report_xgbc.iloc[3,2]}
```

```
[139]: df_report_rf
```

```
[139]:
```

	precision	recall	f1-score	support
0	0.890196	0.886719	0.888454	256.000000
1	0.978770	0.979487	0.979129	1365.000000
accuracy	0.964837	0.964837	0.964837	0.964837
macro avg	0.934483	0.933103	0.933791	1621.000000
weighted avg	0.964782	0.964837	0.964809	1621.000000

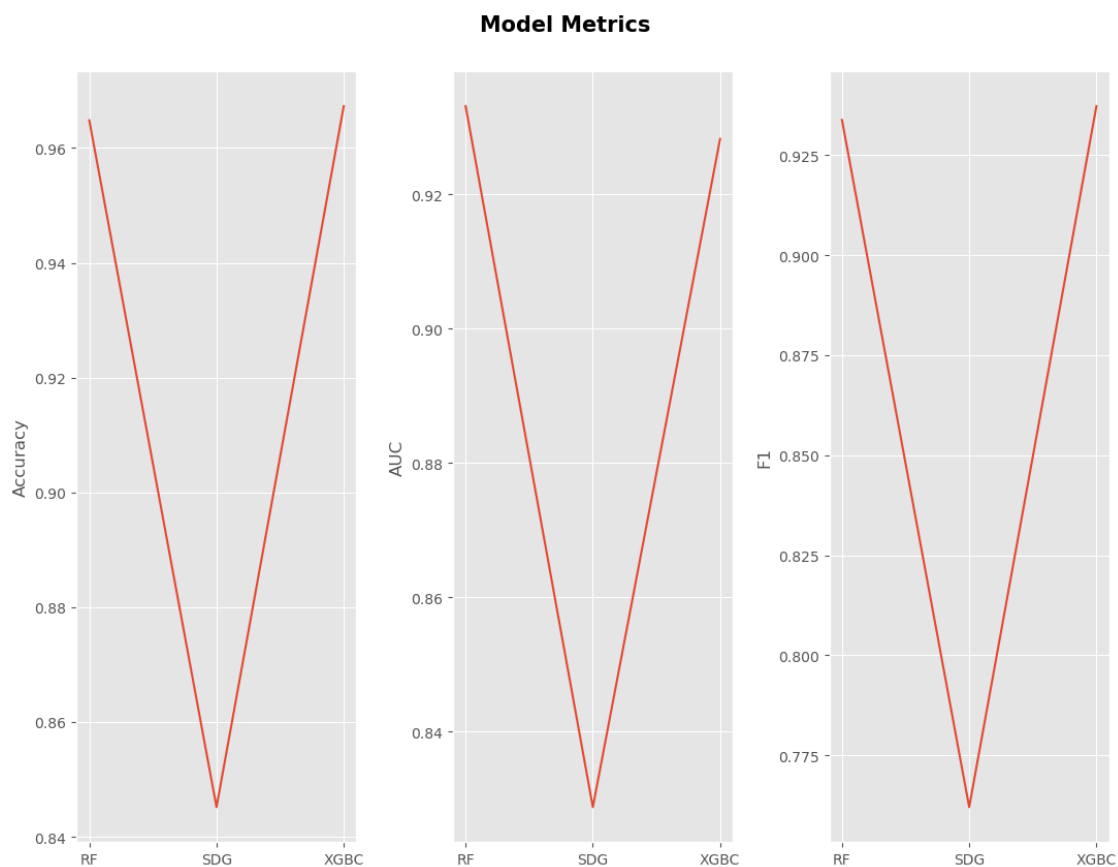
```
[140]: metrics_models = pd.
      ↪ DataFrame([Accuracy,AUC,F1_score],index=["Accuracy","AUC","F1"]).T
      metrics_models
```

```
[140]:      Accuracy      AUC      F1
      RF      0.964837  0.933103  0.933791
      SDG      0.845157  0.828717  0.762047
      XGBC      0.967304  0.928221  0.937233
```

```
[141]: fig,axes= plt.subplots(1,3,figsize=(11,8))

      for i,col in enumerate(metrics_models.columns):
          sns.lineplot(data=metrics_models[col],ax=axes[i] )


      plt.tight_layout()
      plt.suptitle("Model Metrics", fontsize = 15, fontweight = "bold", y=1.05)
      plt.show()
```



```
[120]: false_Negative = [conf_matrix_rf[1,0],
      conf_matrix_sdg[1,0],conf_matrix_xgbc[1,0]]
```

```
false_Negative_S = pd.Series(false_Negative,name="FN",index=["RF","SDG","XGBC"])
```

```
[147]: fig = plt.figure(figsize=(8, 6))
sns.barplot(x=false_Negative_S.index,y=false_Negative_S.values)
plt.ylabel("Number of False Negatives")
plt.title("False Negatives")
plt.show()
```



The RF and XGB model have almost the same f1, with the XGB being slightly higher, but has a greater reduction of false negatives. XGB will be chosen because of its slight increase in f1 but undoubtedly the reduction in false negatives is something of great importance as it represents customers that the model classifies as satisfied with the bank's service but are really customers who are likely to cease using the bank's services.

10 Test

```
[122]: X_Test = test.copy().drop(columns='CLIENTNUM')
```

```
[123]: test_cat_cols = X_Test.select_dtypes(object).columns
test_cat_cols
```

```
[123]: Index(['Gender', 'Education_Level', 'Marital_Status', 'Income_Category',  
          'Card_Category'],  
          dtype='object')
```

```
[124]: X_Test[test_cat_cols] = enc.transform(X_Test[test_cat_cols])
```

```
[125]: X_Test.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 2026 entries, 0 to 2025  
Data columns (total 19 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   Customer_Age                          2026 non-null   int64  
1   Gender                                2026 non-null   float64  
2   Dependent_count                       2026 non-null   int64  
3   Education_Level                       2026 non-null   float64  
4   Marital_Status                       2026 non-null   float64  
5   Income_Category                      2026 non-null   float64  
6   Card_Category                        2026 non-null   float64  
7   Months_on_book                       2026 non-null   int64  
8   Total_Relationship_Count             2026 non-null   int64  
9   Months_Inactive_12_mon              2026 non-null   int64  
10  Contacts_Count_12_mon               2026 non-null   int64  
11  Credit_Limit                        2026 non-null   float64  
12  Total_Revolving_Bal                 2026 non-null   int64  
13  Avg_Open_To_Buy                     2026 non-null   float64  
14  Total_Amt_Chng_Q4_Q1                2026 non-null   float64  
15  Total_Trans_Amt                     2026 non-null   int64  
16  Total_Trans_Ct                      2026 non-null   int64  
17  Total_Ct_Chng_Q4_Q1                 2026 non-null   float64  
18  Avg_Utilization_Ratio                2026 non-null   float64  
dtypes: float64(10), int64(9)  
memory usage: 316.6 KB
```

```
[126]: X_Test[X_Test.columns] = scaler.transform(X_Test)
```

```
[127]: X_Test = X_Test.drop(columns=drop_cols)
```

```
[144]: final_model_prediction = gridXgbc.predict(X_Test)
```

```
[145]: #create Series with predictions  
final_prediction_series = pd.Series(data = final_model_prediction,   
    name='target')  
final_prediction_series
```

```
[145]: 0      1
      1      1
      2      1
      3      1
      4      1
      ..
      2021    1
      2022    0
      2023    1
      2024    1
      2025    1
      Name: target, Length: 2026, dtype: int64
```

```
[146]: #convert series to json file
      json_dict = {'target': final_prediction_series.to_dict()}
      with open('predictions.json', 'w') as f:
          json.dump(json_dict, f)
```